

# Manifold-Constrained Hyper-Connections (mHC): Architecture and Theory

**Manifold-Constrained Hyper-Connections (mHC)** is a recent extension of the standard Transformer residual block designed to widen the residual path while preserving the stability of identity mappings. In a conventional Transformer, each layer computes

$$x_\ell = F(x_{\ell-1}) + x_{\ell-1},$$

so that the input  $x_{\ell-1}$  is passed directly to the output. This *identity mapping* ensures stable training in very deep networks ① ②. In contrast, **Hyper-Connections (HC)** [Zhu et al. 2024] expand the residual stream into  $k$  parallel channels, introducing learnable weights to mix information across these channels.

Specifically, HC replaces the simple skip-connection by three learnable matrices  $W_r$ ,  $W_{\text{out}}$ , and  $W_{\text{in}}$ :  $W_r \in \mathbb{R}^{k \times k}$  mixes the  $k$  residual streams,  $W_{\text{out}} \in \mathbb{R}^k$  collapses them back into a  $d$ -dimensional output, and  $W_{\text{in}} \in \mathbb{R}^{d \times k}$  writes the output back into the streams for the next layer ③. This greatly increases the *information bandwidth* of the residual path, but it also breaks the strict identity property, leading to instability at scale ① ③.

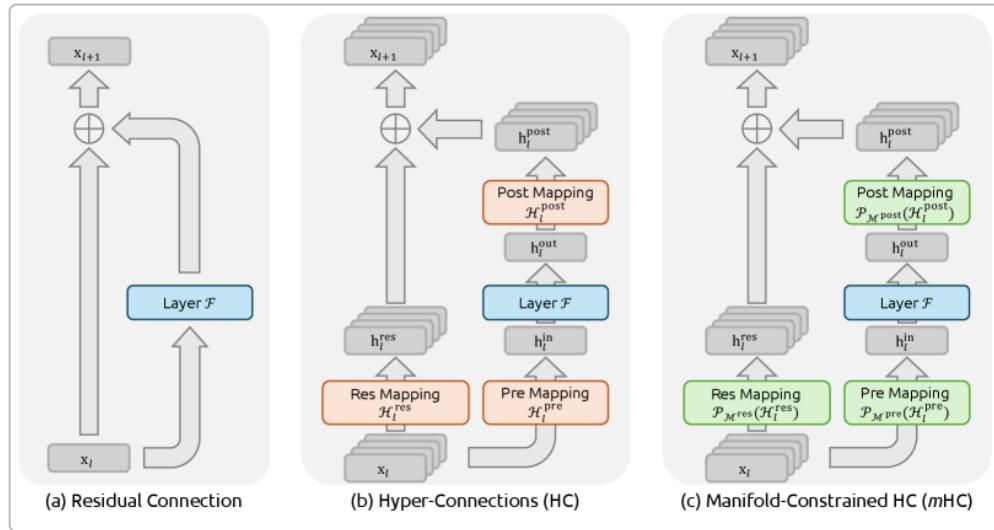


Figure: Illustration of residual paradigms. (a) Standard residual block with identity skip. (b) Hyper-Connections (HC) block with  $k$ -channel residual streams (orange) and learnable mix weights  $W_r, W_{\text{out}}, W_{\text{in}}$ . (c) Manifold-Constrained Hyper-Connections (mHC) enforces  $W_r$  to lie on the doubly-stochastic manifold (green), restoring the mean-preserving identity property while keeping a widened residual stream ① ③.

In mHC, the **core idea** is to constrain the HC mixing matrix  $W_r$  to the *Birkhoff polytope* of doubly stochastic matrices. Concretely, each  $W_r \in \mathbb{R}^{k \times k}$  is enforced to satisfy

$$W_r \mathbf{1} = \mathbf{1}, \quad W_r^T \mathbf{1} = \mathbf{1}, \quad W_{r,ij} \geq 0,$$

so that every row and column sums to 1 <sup>4</sup>. This constraint makes  $W_r$  a convex combination of permutation matrices. Equivalently, multiplying by  $W_r$  performs a weighted averaging of the  $k$  streams. By construction, this **conserves the feature mean** (the average across streams) and bounds the spectral norm to at most 1 <sup>1</sup> <sup>2</sup>. In other words, each stream's signal is non-expansive and the *composition* of such mappings across layers remains identity-like. Indeed, because the set of doubly-stochastic matrices is *closed under multiplication*, the product  $W_r^{(\ell_1)} W_r^{(\ell_2)} \dots$  stays doubly-stochastic, preserving stability throughout the network depth <sup>1</sup> <sup>2</sup>.

Thus, **mHC restores the identity mapping property** by turning the multi-stream residual connection into a *stable feature fusion* layer. The nonnegative, doubly-stochastic  $W_r$  acts as a bounded mixing: no signal can explode or vanish across many layers (its spectral norm is  $\leq 1$ ) <sup>1</sup>. Meanwhile, the residual stream is still  $k$ -times wider, so the model can learn richer cross-layer interactions. The overall architecture of an mHC-enhanced Transformer layer is identical to the standard Transformer except that every skip-connection uses this  $k$ -stream mechanism. In each layer  $\ell$ , one maintains a hidden state  $H^{(\ell)} \in \mathbb{R}^{k \times d}$  (the residual streams). The layer input  $x_{\ell-1} \in \mathbb{R}^d$  is expanded into  $k$  copies to form  $H_{\text{in}}$ , and the output of the layer's feedforward/attention block is added back into the streams. The update within one block is therefore:

- Mix streams:**  $H' = W_r H_{\text{in}}$ , where  $W_r$  is projected to be doubly-stochastic <sup>1</sup>.

2. **Collapse to output:**  $x_\ell = x_{\ell-1} + W_{\text{out}} H'$ , where  $W_{\text{out}}$  (a nonnegative  $1 \times k$  vector) sums the streams into the  $d$ -dimensional output <sup>3</sup>.

3. **Write back:** Update the streams by  $H_{\text{out}} = H' + W_{\text{in}} y$ , where  $y = F(x_{\ell-1})$  is the layer's main output and  $W_{\text{in}}$  (a  $k \times d$  nonnegative matrix) projects that output into the streams <sup>3</sup>.

In summary, an mHC block behaves like a normal residual block except that its skip-connection is replaced by a **manifold-constrained hyper-connection** involving  $W_r$ ,  $W_{\text{out}}$ ,  $W_{\text{in}}$ . A concrete pseudocode of the forward pass is:

```
def mHC_block(x_prev, H_prev):
    # x_prev: d-dimensional input; H_prev: kxd residual streams
    # 1) Flatten streams to form context vector (for dynamic mapping) and
    compute W_r:
        a_dyn = W1 @ H_prev.flatten() + b1
        a_stat = W2 @ H_prev.flatten() + b2
        W_hat = sigmoid(a_dyn + a_stat)          # ensure positivity
        W_exp = exp(W_hat)                      # make entries positive
    # 2) Sinkhorn-Knopp projection to make W_exp doubly-stochastic:
        W = W_exp
        for t in range(N_iter):
            W = W / W.sum(dim=1, keepdims=True)  # normalize rows
            W = W / W.sum(dim=0, keepdims=True)  # normalize columns
    # 3) Apply residual mixing:
        H_mixed = W @ H_prev                  # kxk times kxd -> kxd
    # 4) Compute main layer output and combine:
        y = F(x_prev)                        # e.g. attention or FFN (dimension
        d)
        x_new = x_prev + W_out @ H_mixed      # collapse streams to d-dim and add
        skip
```

```

# 5) Update streams:
H_new = H_mixed + (W_in @ y).reshape(k,d)
return x_new, H_new

```

Here **Step 2** implements the **Sinkhorn-Knopp** algorithm ⑤ to project the unconstrained matrix  $\exp(W_{\text{hat}})$  onto the doubly-stochastic manifold. Each iteration alternately normalizes all rows and all columns to sum to 1. After enough iterations (the paper uses  $N_{\text{iter}} = 20$ ),  $W$  is (approximately) in the Birkhoff polytope ⑤. In training, this projection is part of the forward pass; gradients flow through the Sinkhorn iterations (the authors implement a custom backward pass) so that the original, un-projected weights are learned end-to-end. No extra loss term is needed – the constraint is enforced directly by the projection.

## Mathematical Properties and Formulation

Formally, letting  $\mathcal{B}_k = \{W \in \mathbb{R}^{k \times k} : W\mathbf{1} = \mathbf{1}, W^T\mathbf{1} = \mathbf{1}, W_{ij} \geq 0\}$  be the Birkhoff polytope, mHC enforces  $W_r^{(\ell)} \in \mathcal{B}_k$  for every layer  $\ell$  ④. Equivalently, each layer's mixed output is a convex combination of its  $k$  input streams. This yields three key benefits ① ② :

- **Spectral-Norm Bound:** Any doubly-stochastic matrix has spectral norm  $\leq 1$  ④. Thus  $W_r$  is a non-expansive map. The hidden signal cannot amplify (nor vanish) exponentially across layers, preventing exploding/vanishing gradients.
- **Identity Closure:** Since  $\mathcal{B}_k$  is closed under multiplication, the product of successive  $W_r$  matrices remains doubly-stochastic ②. Therefore the *global* effect of many layers still acts like a weighted identity mapping on average, preserving the feature mean throughout the network ① ②.
- **Geometric Interpretation:** By Birkhoff-von Neumann, a doubly-stochastic  $W$  is a convex mixture of permutation matrices ④. mHC thus continuously mixes streams in a balanced way. Over many layers, information from all streams gets gradually fused monotonically.

Importantly, the input-output identity is recovered as a special case: if  $k = 1$ , the only “matrix” is [1] which exactly matches the original skip-connection. In general  $k > 1$  allows richer mixing, but the conservation of row/column sums ensures that “on average” the identity path remains intact ①. The input and output mappings  $W_{\text{in}}$  and  $W_{\text{out}}$  are also constrained to be nonnegative, avoiding any cancellation across streams. These combined manifold constraints guarantee that mHC behaves like a *well-conditioned* generalization of residual connection.

## Parameterization and Learning

In practice, each weight matrix is implemented using a *dynamic + static* decomposition, as in the original Hyper-Connections formulation ⑥. Concretely, at layer  $\ell$  we flatten the hidden streams  $H^{(\ell-1)}$  into a vector and compute two sets of coefficients via small linear projections: a dynamic component (input-dependent) and a static bias. A sigmoid nonlinearity then enforces positivity. For example, one sets

$$W_{\text{hat}} = \sigma(W_a \text{vec}(H) + W_b \text{vec}(H) + b),$$

where  $\sigma$  is sigmoid and  $b$  is a learnable bias. Learnable **gating scalars**  $\alpha, \beta$  (initialized small) can weight the dynamic vs. static parts ⑥. This yields an unconstrained positive matrix  $W_{\text{exp}} = \exp(W_{\text{hat}})$ . We then perform the Sinkhorn normalization: for iteration  $t = 0, \dots, n - 1$ ,

$$W^{(t+1)} = C(R(W^{(t)})), \quad R(M)_{ij} = \frac{M_{ij}}{\sum_j M_{ij}}, \quad C(M)_{ij} = \frac{M_{ij}}{\sum_i M_{ij}},$$

alternating row and column normalization. As  $t \rightarrow \infty$ ,  $W^{(t)} \rightarrow W_r$  becomes doubly-stochastic <sup>5</sup>. In mHC, the authors fix  $n = 20$  to get an approximate projection <sup>5</sup>. This  $W_r$  is then used to mix streams. Throughout training, gradients are back-propagated through all these steps (the Sinkhorn steps are differentiable with a custom backward pass) so that the underlying parameters  $W_a, W_b, b$  are learned to produce stable  $W_r$ .

The other mappings  $W_{\text{out}} \in \mathbb{R}^{1 \times k}$  and  $W_{\text{in}} \in \mathbb{R}^{k \times 1}$  are learned similarly but without a global normalization: typically one ensures these are nonnegative (e.g. by ReLU or Sigmoid) so that adding and writing into streams does not cancel signals. In ablations, the authors found  $W_r$  (the residual mixing) was the most important of the three mappings <sup>6</sup>. No additional loss terms are used: the model is trained end-to-end on the usual language-modeling objective (e.g. cross-entropy over tokens). The manifold constraint itself is enforced solely by the above parameterization and Sinkhorn projection.

## Training Stability and Efficiency

Empirically, mHC enables stable training of very large models. In large-scale experiments (up to 27B parameters), unconstrained HC exhibited catastrophic loss spikes and gradient explosions. In contrast, mHC-trained models show smooth loss curves and bounded gradient norms <sup>7</sup> <sup>1</sup>. In other words, the manifold constraint effectively restores the “identity-map conservation” of residual nets. The loss gap between HC and mHC grows as models scale, confirming that mHC removes a structural instability <sup>7</sup>. Across multiple benchmarks, mHC models consistently outperformed standard Transformers of the same size, indicating the gains are architectural, not just hyperparameter tweaks <sup>1</sup> <sup>7</sup>.

Of course, widening the residual stream increases memory I/O. To make mHC practical, DeepSeek engineers applied several optimizations <sup>8</sup> <sup>9</sup>. Key strategies include **kernel fusion** (combining the normalization, projection and residual merge into efficient GPU kernels), **selective recomputation** (discarding and later recomputing intermediate activations for the streams), and overlapping communication in pipeline parallelism. These optimizations limit the added overhead: in reported runs a 4× wider stream incurred only ~6–7% extra training time <sup>8</sup>. The result is that mHC can be deployed at scale with reasonable cost, unlocking a higher-capacity residual topology without the usual training collapse.

## Related Concepts

- **Residual and Highway Networks:** mHC can be seen as a generalization of Highway/Residual designs <sup>3</sup>. Highway networks introduced gated skip connections, and later designs like DenseNet or Deep Layer Aggregation added cross-layer links. HC/mHC instead add *wide, weighted* skip connections.
- **Hypernetworks:** Though similarly named, traditional *Hypernetworks* (Ha et al. 2016) are different: they use one neural network to **generate weights** for another <sup>10</sup>. mHC’s “hyper-connections” do not generate other weights; rather, they **mix** the residual streams with trainable matrices.
- **Manifold Constraints in Learning:** Constraining weights to manifolds is a known technique. For instance, enforcing orthonormal weight matrices (on the Stiefel manifold) ensures spectral norm =1 in layers, which also improves stability. mHC’s use of the *doubly stochastic* manifold is a novel

instance of this principle [1](#) [11](#). By projecting onto a Riemannian manifold with known closure properties, mHC provides rigorous guarantees about training dynamics.

- **Related Works:** Beyond HC, other “multi-stream” architectures have been proposed (e.g. Residual Matrix Transformer, MUDDFormer, etc. cited in [12](#) [11](#)), but most lacked a manifold constraint and therefore hit scalability issues. The mHC-GNN work [11](#) [2](#) adapts these ideas to graph networks, and theoretical analyses there mirror mHC’s claim: constraining to a doubly-stochastic mix *prevents feature collapse* in deep message passing, just as it prevents gradient explosion here.

## End-to-End Walkthrough

Putting it all together, the *end-to-end* forward pass of an mHC-enhanced Transformer block is:

1. **Pre-Normalize** the input  $x_{\ell-1}$  (as in a standard Transformer).
2. **Compute Residual Mixing:** Expand  $x_{\ell-1}$  into  $k$  streams to form a hidden matrix  $H^{(\ell-1)}$ . Compute the combined coefficients (dynamic + static) and apply the Sinkhorn projection to get  $W_r^{(\ell)} \in \mathcal{B}_k$  [13](#) [1](#).
3. **Mix Streams:** Multiply  $H^{(\ell-1)}$  by  $W_r^{(\ell)}$  to get  $H_{\text{mixed}}^{(\ell)}$ .
4. **Feedforward/Attention:** Process  $x_{\ell-1}$  through the attention or feedforward sublayer to get  $y_\ell$ .
5. **Merge Outputs:** Collapse the mixed streams via  $W_{\text{out}}^{(\ell)}$  and add to the input:  $x_\ell = x_{\ell-1} + W_{\text{out}}^{(\ell)} H_{\text{mixed}}^{(\ell)}$ .
6. **Update Streams:** Write the new output into the streams via  $W_{\text{in}}^{(\ell)}$ : set  $H_{\text{new}}^{(\ell)} = H_{\text{mixed}}^{(\ell)} + W_{\text{in}}^{(\ell)} y_\ell$ .
7. **Proceed:** Pass  $x_\ell$  to the next layer (and  $H_{\text{new}}^{(\ell)}$  as the updated streams).

Throughout, all weight parameters  $\{W_a, W_b, b, \alpha, \beta\}$  that define  $W_r$ , as well as  $W_{\text{in}}, W_{\text{out}}$ , are learned by gradient descent on the usual language-model objective. The Sinkhorn projection is simply part of the forward graph. In this way, mHC ties together theory and practice: the **theoretical requirement** (rows/columns sum to 1) is enforced algorithmically at each step, and the **implementation** (kernels, fused operations) ensures it runs efficiently.

In summary, Manifold-Constrained Hyper-Connections widen the residual pathway (like HC) but **rigidly constrain** the connections to a stability-preserving manifold (via Sinkhorn) [1](#) [11](#). This yields a new class of Transformer layers that combine high capacity with provable training stability, suggesting a promising path for next-generation deep networks.

**Sources:** The above explanation is based on the *mHC* paper by Xie et al. (2026) [1](#) [13](#) and related references [3](#) [11](#) [10](#). All equations and figures are adapted from these sources.

[1](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) [9](#) [12](#) [13](#) [2512.24880] mHC: Manifold-Constrained Hyper-Connections

<https://arxiv.labs.arxiv.org/html/2512.24880>

[2](#) [11](#) mHC-GNN: Manifold-Constrained Hyper-Connections for Graph Neural Networks

<https://arxiv.org/html/2601.02451v1>

[10](#) [1609.09106] HyperNetworks

<https://arxiv.org/abs/1609.09106>